

Intro to TensorFlow

Ben Kubit

NSC211 Spring 2017

University of California, Davis

Overview

- TF Python API
- Building a network
- XOR example
- TensorBoard

TF Python API

- Graphs and Sessions:
 - A **computational graph** is a series of TensorFlow operations arranged into a graph of nodes.
 - To actually evaluate the nodes, we must run the computational graph within a **session**.
 - Graph is run entirely outside Python. TF relies on efficient C++ backend for computations.
 - Connection to this backend is called a **session**.
 - Since TF knows the entire **computation graph**, faster computing
 - E.g., automatic differentiation to find loss gradients with respect to each variable.

TF Python API

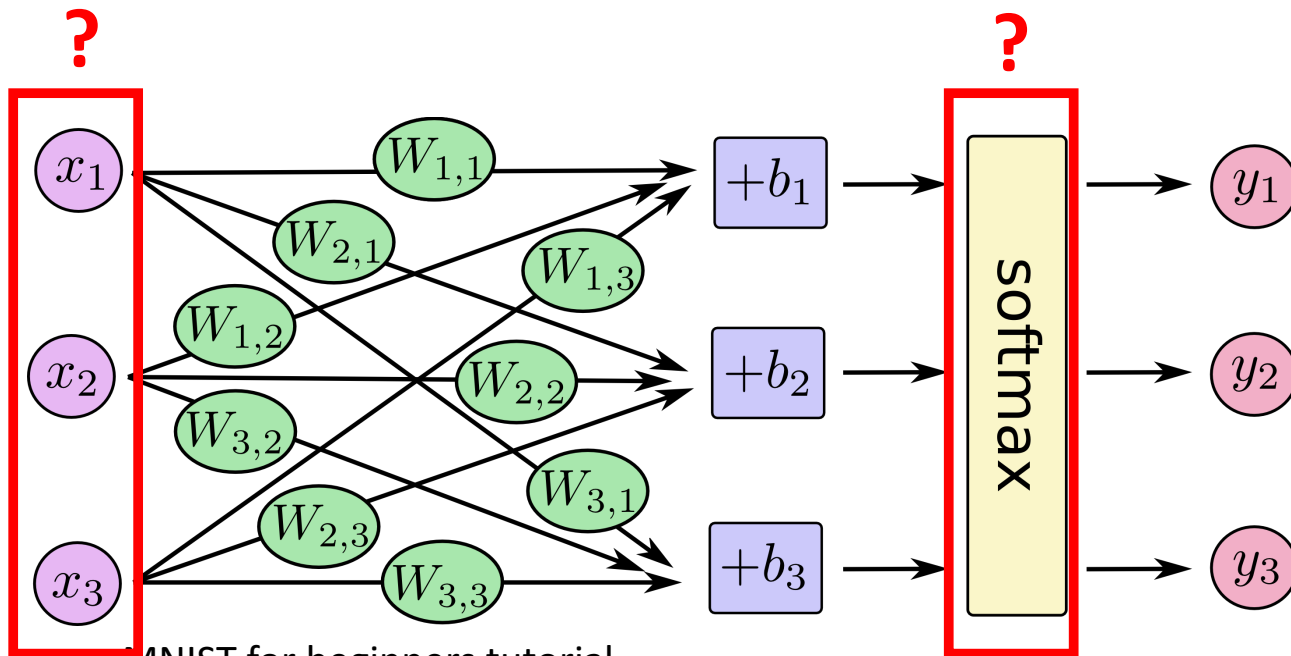
- *Tensor Nodes*: may be thought of as an operation that takes no inputs and always produces the same output corresponding to the constant/variable it represents.
 - **Constants**: hard-coded values.
 - **Placeholders**: promise to provide a value later. Specify **placeholder** values when you run a session (typically input data/predicted values).
 - **Variables**: allow us to add trainable parameters to a graph. Initialize on first run, then dynamically update based on optimization function (typically model parameters).
- *Operation Nodes*: express the combination/transformation of data flowing through the graph
 - E.g., `add_node = y + x` -> receives input from tensor nodes `y` and `x`
 - E.g., `lm_node = tf.matmul(x, w) + b`
- *Summary Nodes*: capture information during training for use with TensorBoard

Building a Network

- Inference - Build the graph
 - Initialize operation nodes and tensor nodes
 - Form of output units \rightarrow functions + # of nodes
- Loss - Choose loss function
 - E.g., MSE, cross-entropy
- Training - Choose optimizer
 - E.g., Gradient descent

Building a Network


- What is a layer?!
 - Project inputs into a new space.
 - Has weights for each incoming edge and biases for each node in the layer.



MNIST for beginners tutorial

https://www.tensorflow.org/get_started/mnist/beginners

Let's try 5 fully-connected layers !



Tensorflow and deep learning - without a PhD by Martin Görner
<https://www.youtube.com/watch?v=vq2nnJ4g6N0>

Deep Learning 6.1 Example: Learning XOR

- The XOR function (“exclusive or”):
 - operation on two binary values, x_1 and x_2 .
 - when only one of these values==1, the function returns 1, otherwise 0
 - 2-way classification
 - We can treat this problem as a regression problem
 - We want our network to perform correctly on the four points:
 - $X = \{[0,0], [0,1], [1,0], \text{ and } [1,1]\}$.
- Examine:
 - utility of hidden layer transformations
 - Influence of initial weight values
 - Low-level ML (vs `tf.contrib.train` high-level ML)
- Code:
 - `NSC211_BKlecture_code.ipynb`

Deep Learning 6.1 Example: Learning XOR

- We can minimize in closed form with respect to w and b using the normal equations.
 - After solving the normal equations, we obtain $w = 0$ and $b = 1/2$.
 - The linear 2 model simply outputs 0.5 everywhere!

```
sess = tf.Session()
#input data
XOR_X = [[0,0],[0,1],[1,0],[1,1]] #input
#XOR_Y = [[0],[1],[1],[0]] #predicted

#placeholders
x_ = tf.placeholder(tf.float32, shape=[4,2], name="x-input")
#use weights/biases from book example
w = tf.Variable(tf.zeros([2,1]), tf.float32)
b = tf.Variable([1/2.], tf.float32)

init = tf.global_variables_initializer()
sess.run(init)

#operation node
linear_model = tf.matmul(x_,w) + b

#see what the predictions are
print(sess.run(linear_model, {x_: XOR_X}))
```




Output:

[[0.5]]
[0.5]		
[0.5]		
[0.5]		


Single fully-connected layer: 2 inputs (so 2 weights total); one bias for the one node

Deep Learning 6.1 Example: Learning XOR

Network 1



Network 2



Deep Learning 6.1 Example: Learning XOR

- solve using a model that learns a different feature space in which a linear model is able to represent the desired solution.
 - simple feedforward network with one hidden layer -> change what is given to output layer

```
sess2 = tf.Session()
#input data
XOR_X = [[0,0],[0,1],[1,0],[1,1]] #input
#XOR_Y = [[0],[1],[1],[0]] #predicted

#placeholders
x_ = tf.placeholder(tf.float32, shape=[4,2], name="x-input")
#use weights/biases from book example
w1 = tf.Variable(tf.ones([2,2]), tf.float32) #W
w2 = tf.Variable([[1.],[-2.]], tf.float32) #w
b1 = tf.Variable([[0.,-1.]], tf.float32) #c
b2 = tf.Variable(tf.zeros(1), tf.float32) #b

init2 = tf.global_variables_initializer()
sess2.run(init2)

#operation nodes
transformedH = tf.nn.relu(tf.matmul(x_,w1) + b1, name=None) #hidden layer with rect. linear act. func.
linear_model = tf.matmul(transformedH, w2) + b2

#see what the predictions are
print(sess2.run(linear_model, {x_: XOR_X}))
```

Output:

[0.]]
[1.]]
[1.]]
[0.]]

1st fully-connected layer: 4 inputs (so 4 weights total); 2 bias for the 2 nodes

2nd fully-connected layer: 2 inputs (so 2 weights total); 1 bias for the single node

Deep Learning 6.1 Example: Learning XOR

- In a real situation, there are lots of model parameters and training examples
 - we cannot guess the solution as we did before.
- now solve the same problem but let's use gradient-based optimization to find params
 - in order to do so need to measure error/cost
 - also need predicted values!

```
sess3 = tf.Session()
#input data
XOR_X = [[0,0],[0,1],[1,0],[1,1]] #input
XOR_Y = [[0],[1],[1],[0]] #predicted

#placeholders, now we need one for predicted vals too
x_ = tf.placeholder(tf.float32, shape=[4,2], name="x-input")
y_ = tf.placeholder(tf.float32, shape=[4,1], name="y-input")
#now we will init with some random values
w1 = tf.Variable(tf.random_uniform([2,2], -2, 2), tf.float32) #W
w2 = tf.Variable(tf.random_uniform([2,1], -2, 2), tf.float32) #w
b1 = tf.Variable(tf.zeros([2]), tf.float32) #c
b2 = tf.Variable(tf.zeros([1]), tf.float32) #b
```

Predicted values (labels)

Initialize variables with random values and zeros

Deep Learning 6.1 Example: Learning XOR

```
#operation nodes
transformedH = tf.nn.relu(tf.matmul(x_,w1) + b1) #hidden layer with rect. linear act. func.
linear_model = tf.matmul(transformedH, w2) + b2

#MSE
loss = tf.reduce_sum(tf.square(linear_model - y_)) #create error vector.We call tf.square to square that error.

#gradient descent
optimizer = tf.train.GradientDescentOptimizer(0.01) #0.01 is learning rate
train = optimizer.minimize(loss) #feed optimizer loss function

init3 = tf.global_variables_initializer()
sess3.run(init3)

#train it
for i in range(10000):
    sess3.run(train, {x_: XOR_X, y_: XOR_Y})

#take a look at the results
predictions = sess3.run(linear_model, {x_: XOR_X})
curr_w1, curr_w2, curr_b1, curr_b2, curr_loss = sess3.run([w1, w2, b1, b2, loss], {x_: XOR_X, y_: XOR_Y})
hidlay = sess3.run(transformedH, {x_: XOR_X, y_: XOR_Y})
print("predictions:\n %s\n hlayinput:\n %s\n"%(predictions,hidlay))
print("w1:\n %s \nw2:\n %s \nb1: %s \nb2: %s \nloss: %s"%(curr_w1, curr_w2, curr_b1, curr_b2, curr_loss))
```

Feed 2nd layer output into loss function

Feed loss function output into optimizer

run optimizer (updates param vars)

Output1:	Output2:	Output3:
predictions: [[4.99999642e-01] [9.99998569e-01] [4.99999642e-01] [1.46031380e-06]]	predictions: [[2.81445682e-06] [9.99997795e-01] [9.99997795e-01] [1.61863863e-06]]	predictions: [[5.00001073e-01] [5.00001073e-01] [9.99999583e-01] [-2.38418579e-07]]

Deep Learning 6.1 Example: Learning XOR

- using this approach we often find different solutions because the minima found depends on the rand. initial weights

Output1:

```
predictions:
[[ 4.99999642e-01]
 [ 9.99998569e-01]
 [ 4.99999642e-01]
 [ 1.46031380e-06]]
hlayinput:
[[ 0.          0.          ]
 [ 0.81953818  1.01845491]
 [ 0.          0.          ]
 [ 0.          1.48590386]]

w1:
[[-0.8780849  0.46744898]
 [ 1.40672278  1.49006176]]
w2:
[[ 1.02826595]
 [-0.3364943 ]]
b1: [-0.58718461 -0.47160682]
b2: [ 0.49999964]
loss: 0.5
```

Output2:

```
predictions:
[[ 2.81445682e-06]
 [ 9.99997795e-01]
 [ 9.99997795e-01]
 [ 1.61863863e-06]]
hlayinput:
[[ 0.          0.00977176]
 [ 0.          1.08689225]
 [ 0.          1.08689225]
 [ 1.17514503  2.16401291]]

w1:
[[ 1.17514503  1.07712054]
 [ 1.17514503  1.07712054]]
w2:
[[-1.70191026]
 [ 0.92839658]]
b1: [-1.17514503  0.00977176]
b2: [-0.00906925]
loss: 2.02685e-11
```

Output3:

```
predictions:
[[ 5.00001073e-01]
 [ 5.00001073e-01]
 [ 9.99999583e-01]
 [-2.38418579e-07]]
hlayinput:
[[ 0.          0.          ]
 [ 0.          0.          ]
 [ 0.13460982  1.74021423]
 [ 0.38378608  0.          ]]

w1:
[[ 1.02023137  1.80393577]
 [ 0.24917631 -1.88829768]]
w2:
[[-1.30281258]
 [ 0.38809583]]
b1: [-0.88562155 -0.06372153]
b2: [ 0.50000107]
loss: 0.5
```


Deep Learning 6.1 Example: Learning XOR

- if we set the weights closer to the values provided in the example the output is more consistent and similar to predicted values

Output:

```
sess4 = tf.Session()
#input data
XOR_X = [[0,0],[0,1],[1,0],[1,1]] #input
XOR_Y = [[0],[1],[1],[0]] #predicted

#placeholders, now we need one for predicted vals too
x_ = tf.placeholder(tf.float32, shape=[4,2], name="x-input")
y_ = tf.placeholder(tf.float32, shape=[4,1], name="y-input")
#constrain rand. values
w1 = tf.Variable(tf.random_uniform([2,2], .7, 1.3), tf.float32) #W
w2 = tf.Variable(tf.random_uniform([2,1], -2, 1), tf.float32) #w
b1 = tf.Variable(tf.zeros([2]), tf.float32) #c
b2 = tf.Variable(tf.zeros([1]), tf.float32) #b
```

Initialize variables with
random values and zeros



-
-
-

W(w1): $\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$

w(w2): $\begin{bmatrix} -2 \\ 1 \end{bmatrix}$

b1(c): $[-1 \ 0]$

b2(b): $[0]$

```
predictions:
[[ 3.11833674e-06]
 [ 9.99997497e-01]
 [ 9.99997616e-01]
 [ 1.76924414e-06]]
hlayinput:
[[ 0.00000000e+00  3.96847035e-08]
 [ 0.00000000e+00  1.05011714e+00]
 [ 0.00000000e+00  1.05011725e+00]
 [ 1.31451106e+00  2.10023451e+00]]

w1:
[[ 1.31451106  1.05011725]
 [ 1.31451106  1.05011714]]
w2:
[[-1.5214709 ]
 [ 0.95226938]]
b1: [ -1.31451106e+00  3.96847035e-08]
b2: [ 3.08054632e-06]
loss: 2.48056e-11
```

Deep Learning 6.1 Example: Learning XOR

- Summary
 - Utility of hidden layer transformations
 - Importance of initial parameter values
- Better way to save diagnostic info?
 - TensorBoard
 - Visualize the graph
 - Generate summary nodes

Better way to save diagnostic info?

- TensorBoard

collect all summary values

Create object to output summaries

Gather summary values for this iter

Write out summaries to file and refresh

```
linear_model = tf.matmul(transformedH, w2) + b2
tf.summary.histogram("predicted", linear_model)

#MSE
loss = tf.reduce_sum(tf.square(linear_model - y_))
tf.summary.scalar("curr_loss", loss)


# Build the summary Tensor based on the TF collection of summaries.
summary = tf.summary.merge_all()
init = tf.global_variables_initializer()
sess = tf.Session()

# Instantiate a SummaryWriter to output summaries and the Graph.
summary_writer = tf.summary.FileWriter(log_dir, sess.graph)
sess.run(init)
#train it
for i in range(10000):
    #run the session and get summary info
    _, suminfo = sess.run([train, summary], feed_dict={x_: XOR_X, y_: XOR_Y})
    # Write the summaries every 100 trials.
    if i % 100 == 0:
        summary_writer.add_summary(suminfo, i) # Update the events file.
        summary_writer.flush()
```

collect prediction values


collect loss value

TensorBoard Summary Scalar

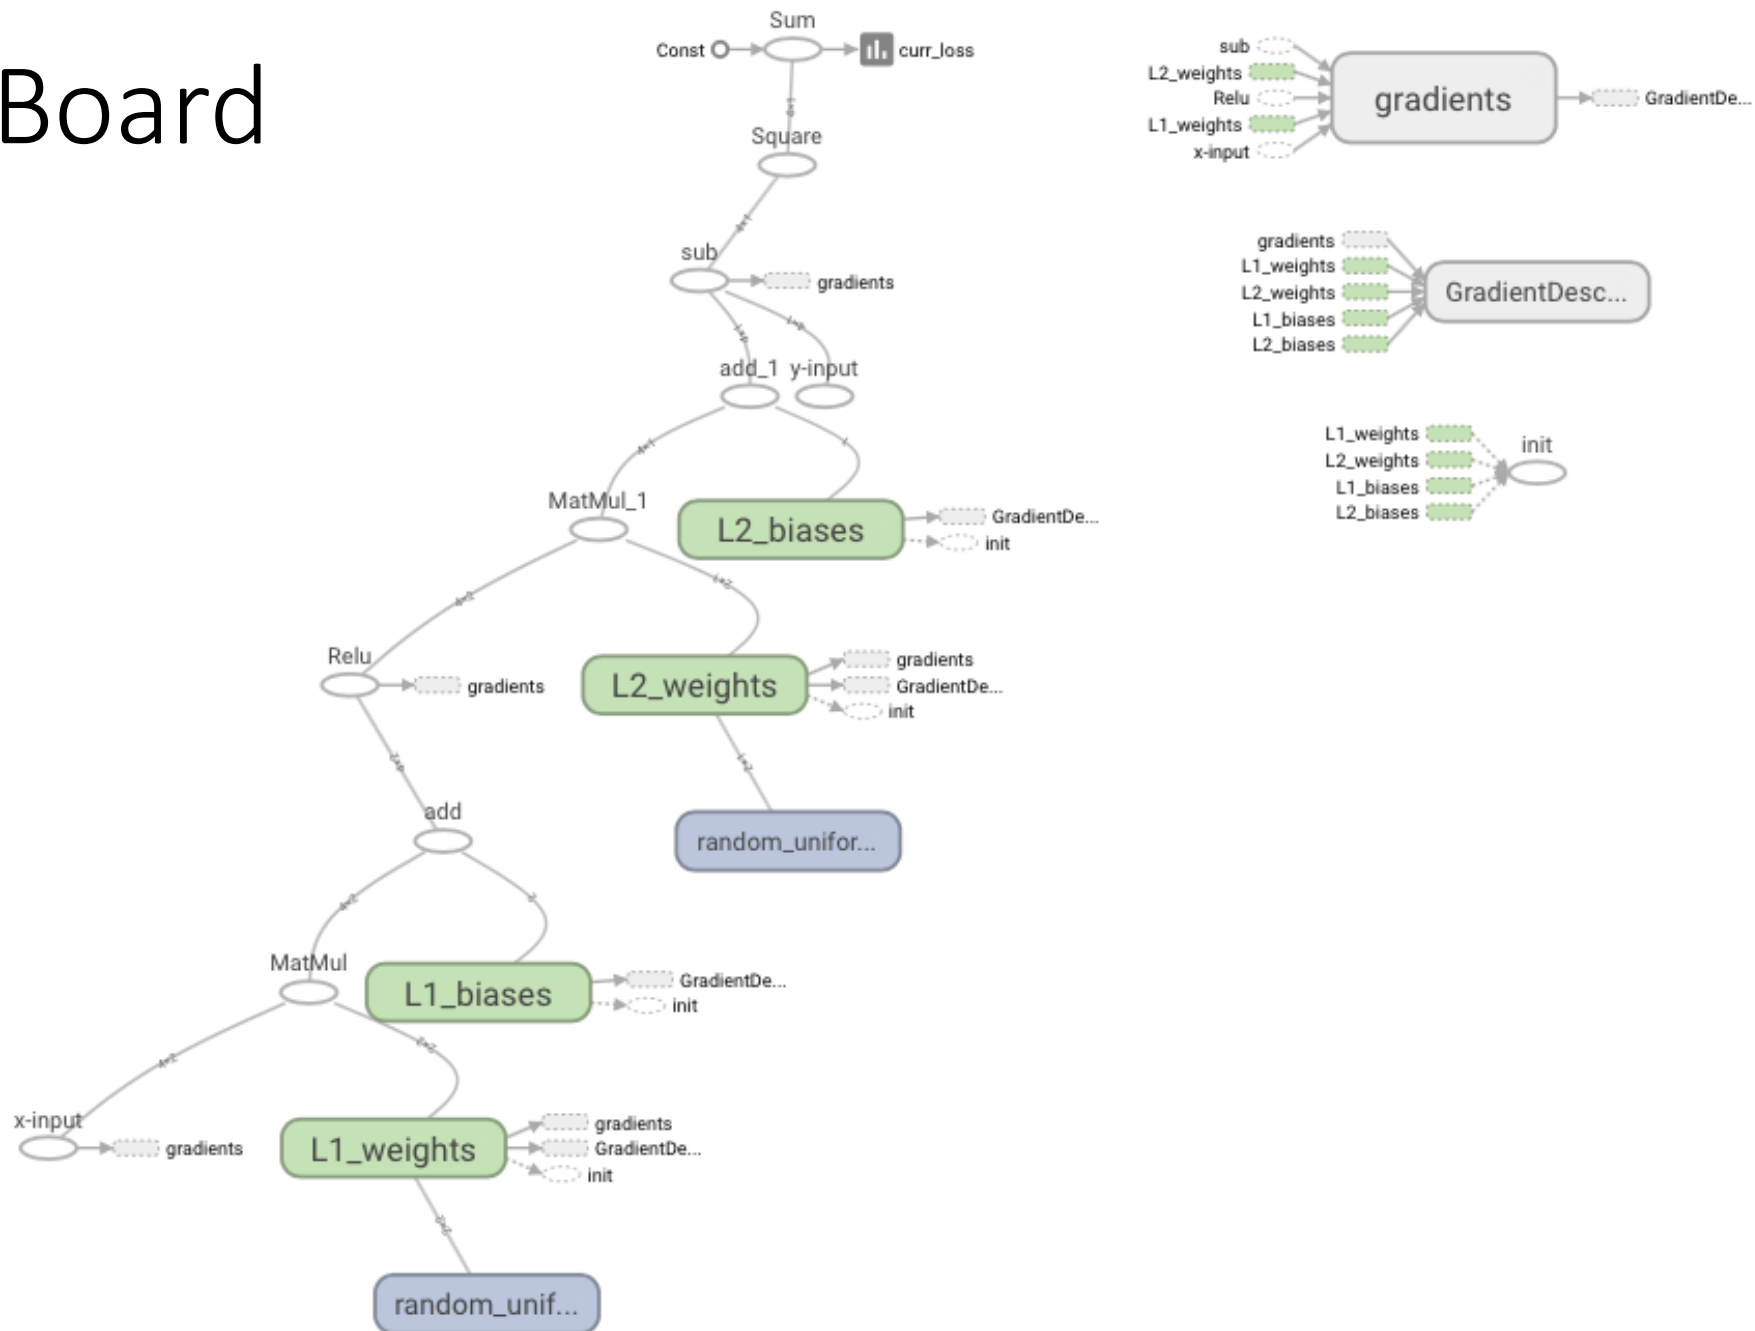


TensorBoard

Summary Histogram



TensorBoard Graph



Summary

- TF Python API
 - https://www.tensorflow.org/get_started/get_started
 - https://www.tensorflow.org/programmers_guide/
 - https://www.tensorflow.org/api_docs/python/
- Building a network
 - <https://www.youtube.com/watch?v=vq2nnJ4g6N0>
- XOR TF example:
 - NSC211_BKlecture_code.ipynb
- Diagnostic tools:
 - https://www.tensorflow.org/get_started/summaries_and_tensorboard
 - https://www.tensorflow.org/api_guides/python/summary
 - https://www.tensorflow.org/api_guides/python/tfdbg
 - https://www.tensorflow.org/api_docs/python/tf/InteractiveSession

Additional considerations

- High-level API
 - E.g., `tf.contrib.train`
- Data preparation
 - E.g., batch normalization
- Batch processing
 - E.g., cross-validation